



Performance Optimization and Best Practices in SQL

Learning Objectives

By the end of this lesson, you will be able to:

- 🕒 Implement an execution plan in SQL
- 🕒 Compare VARCHAR, CHAR, and NVARCHAR
- 🕒 List the index guidelines and analyze the creation of XML PATH and INDEX
- 🕒 Outline queries with SQL indexes



Execution Plan in SQL

Execution Plan in SQL



Execution plan

- Unlike other database solutions, MySQL does not generate byte-code to execute a query result. Instead, the query execution plan is used.
- The query execution plan is a list of instructions that the query execution must follow to produce the query result.
- It converts the source code (SQL query) into an executable program.

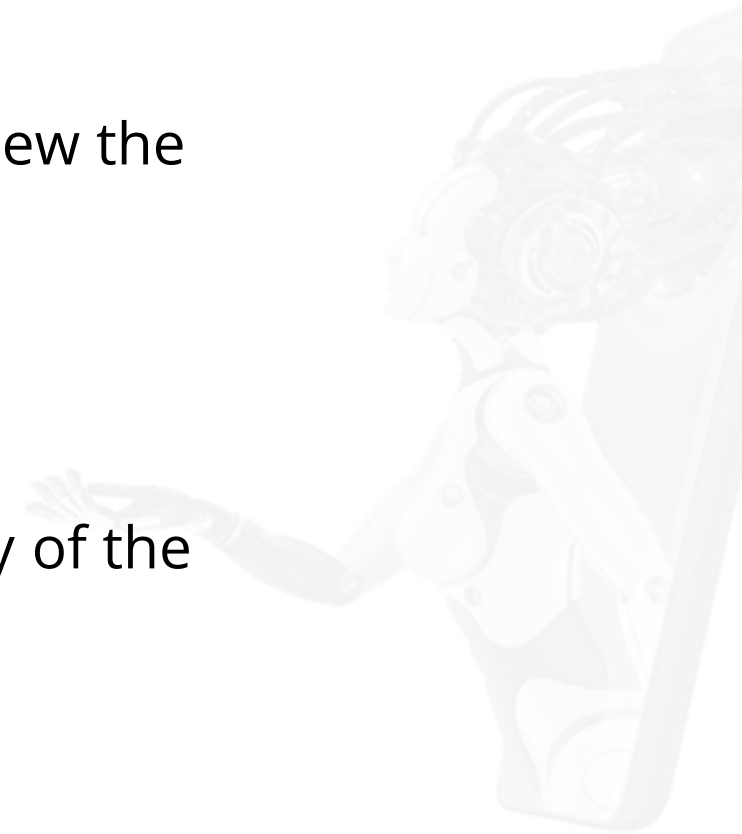
Execution Plan in SQL

Problem Scenario:

A data analyst of a company wants to check the execution plan for a query to view the performance and cost of the query.

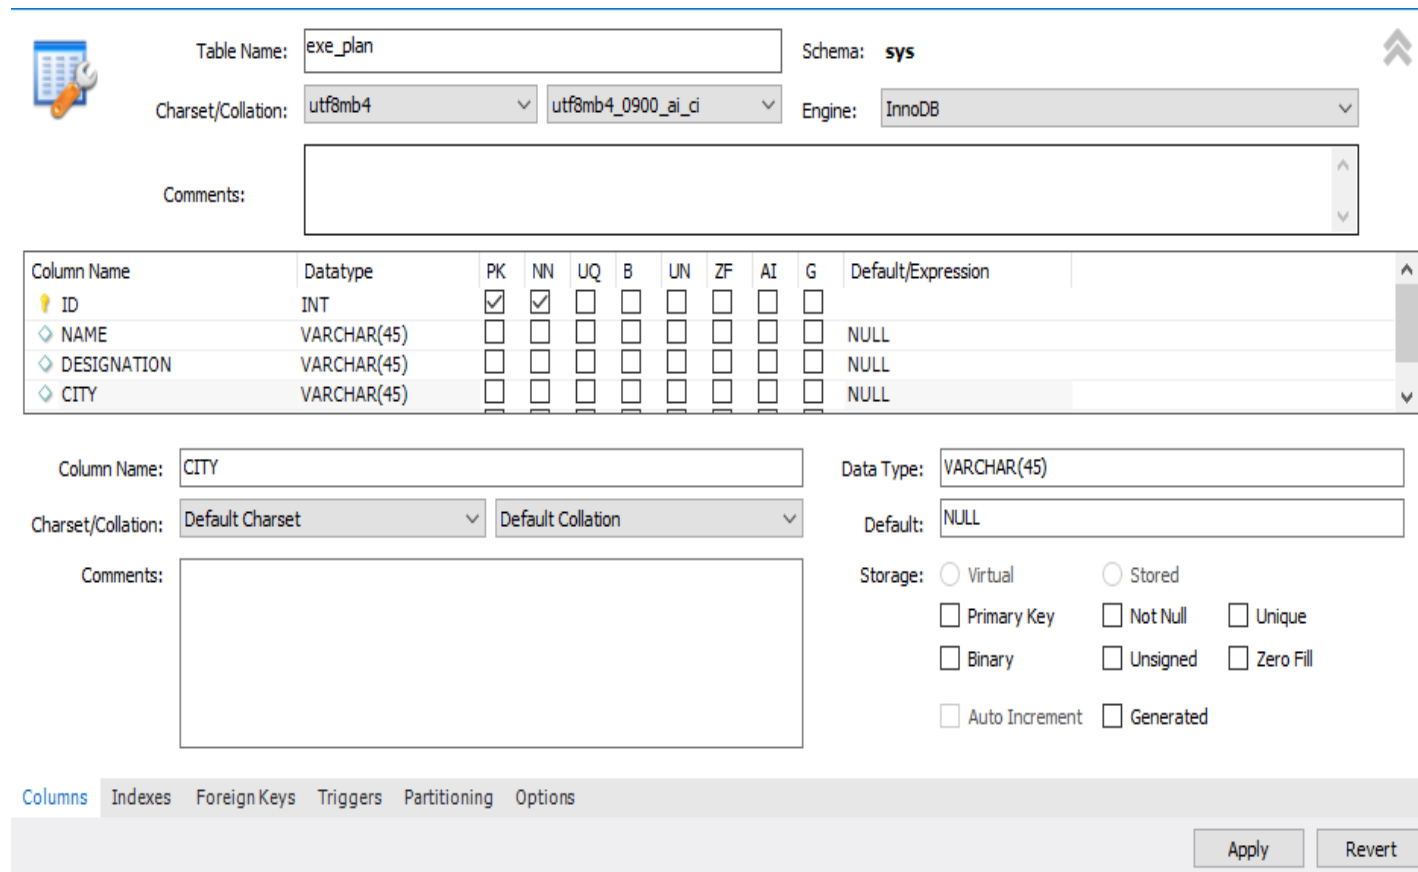
Objective:

View the execution plan and create an index on the query to improve efficiency of the query.



Execution Plan in SQL

Create table Emp_info



The screenshot shows the 'New Table' dialog in SQL Server Enterprise Manager. The table name is 'exe_plan' and the schema is 'sys'. The charset/collation is 'utf8mb4' and the engine is 'InnoDB'. The table has four columns: ID (INT, PK, NN), NAME (VARCHAR(45)), DESIGNATION (VARCHAR(45)), and CITY (VARCHAR(45)). The 'CITY' column is currently selected for editing.

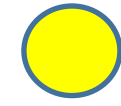
Column Name	Datatype	PK	NN	UQ	B	UN	ZF	AI	G	Default/Expression
ID	INT	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	
NAME	VARCHAR(45)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
DESIGNATION	VARCHAR(45)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
CITY	VARCHAR(45)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL

Column Name: CITY
Data Type: VARCHAR(45)
Charset/Collation: Default Charset
Default: NULL
Storage: ☐ Virtual ☐ Stored
☐ Primary Key ☐ Not Null ☐ Unique
☐ Binary ☐ Unsigned ☐ Zero Fill
☐ Auto Increment ☐ Generated

QUERY

```
CREATE TABLE EXE_PLAN (  
  ID int,  
  NAME varchar(255),  
  DESIGNATION varchar(255),  
  CITY varchar(255),  
);
```


Execution Plan in SQL



Insert data in a table

```
1 • INSERT INTO sys.exe_plan (ID,NAME,DESIGNATION,CITY) VALUES ('4', 'KARENE', 'SLD', 'Bangkok');
2 • select *from sys.exe_plan;
```

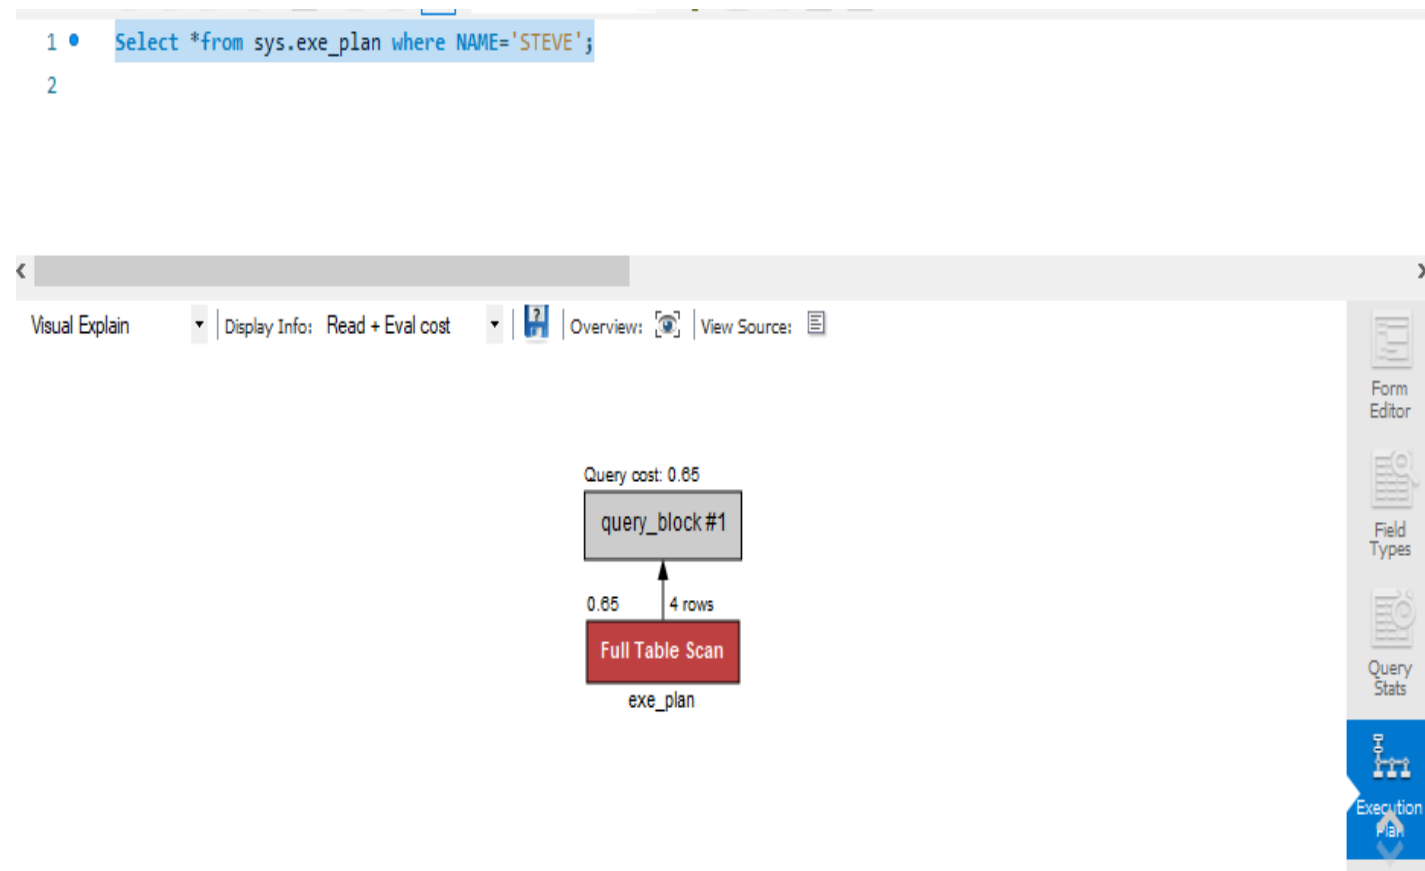
ID	NAME	DESIGNATION	CITY
1	ROY	SLD	Noida
2	KATRINA	LD	New York
3	STEVE	LD	Paris
4	KARENE	SLD	Bangkok
*	NULL	NULL	NULL

QUERY

```
INSERT INTO sys.exe_plan (ID,NAME,DESIGNATION,CITY)
VALUES ('3', 'STEVE', 'LD', 'Paris');
```

```
Select *from sys.exe_plan;
```

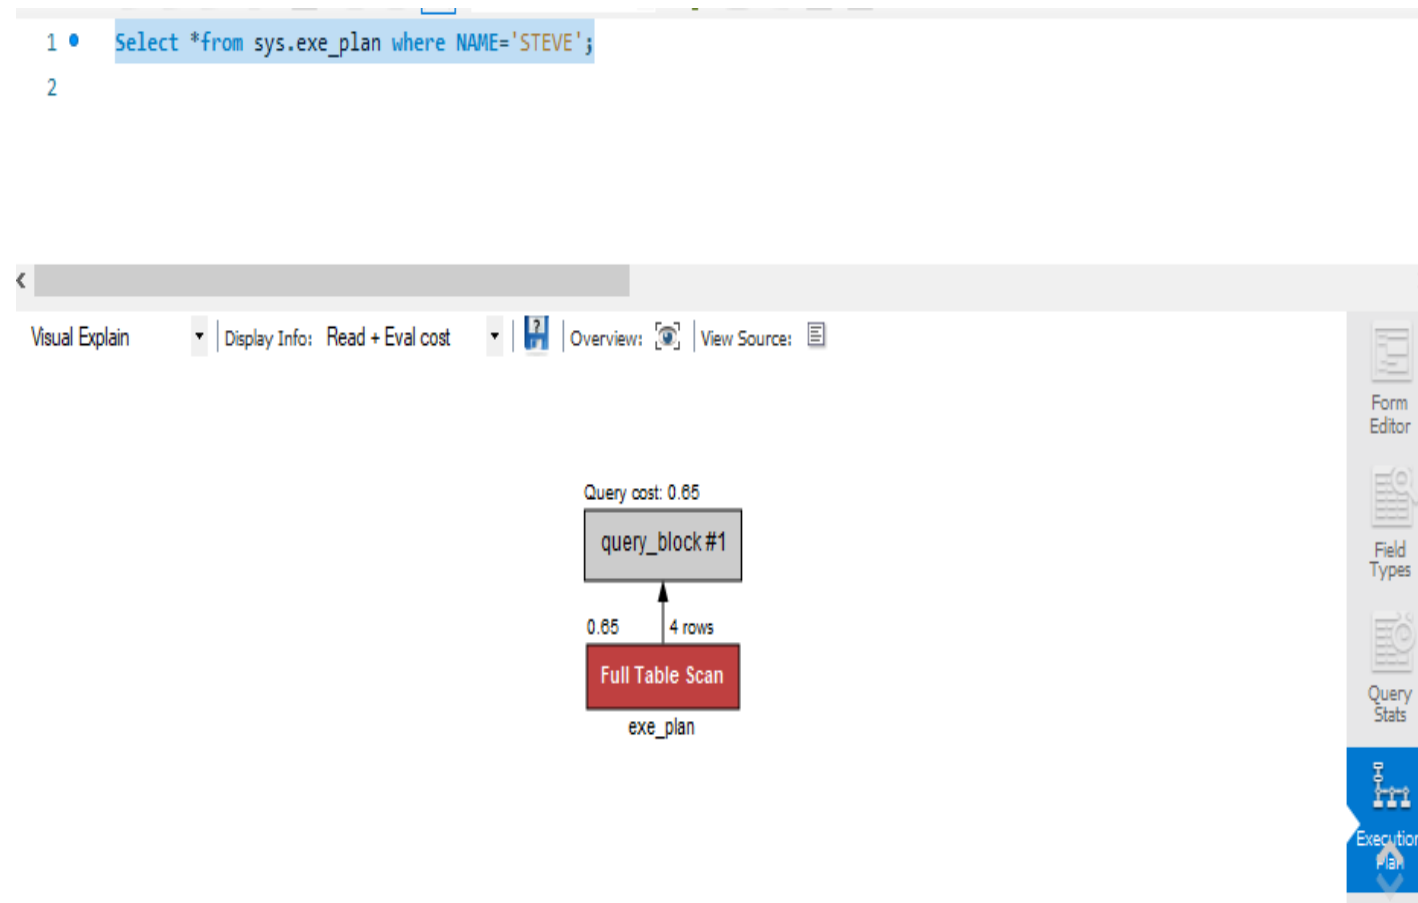
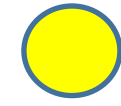
Execution Plan in SQL



QUERY

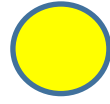
```
select *from sys.exe_plan where NAME='STEVE';
```


Execution Plan in SQL

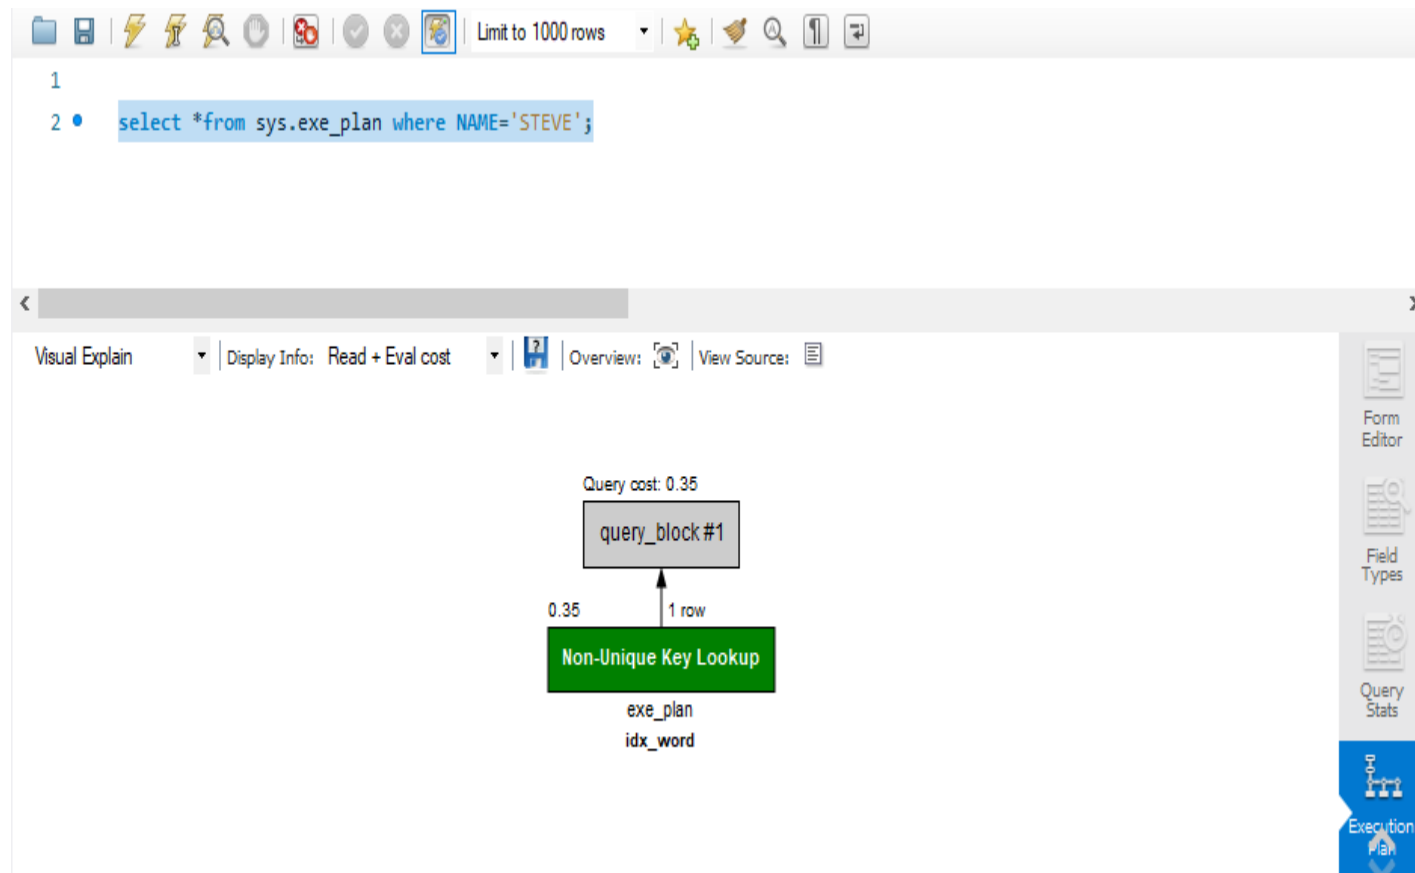


- The image shows the execution plan, and the red box which is in the image is due to the Performance and high cost of the query.
- Query cost is a metric used in MySQL to determine how expensive a query is in terms of the overall cost of query execution.

Execution Plan in SQL



Creating an index to enhance the query performance



QUERY

```
create index idx_word on exe_plan(NAME);  
  
select *from sys.exe_plan where NAME='STEVE';
```

Difference Between CHAR, VARCHAR, and NVARCHAR

Difference Between CHAR, VARCHAR, and NVARCHAR

CHAR

It's a data type with a set length.

Non-Unicode characters are stored here.

Each character is given one byte of space.

VARCHAR

It's a data type with a changeable length.

Non-Unicode characters are stored here.

Each character takes one byte of memory.

NVARCHAR

It's a data type with a changeable length.

Unicode characters are stored here.

Each character takes two bytes of memory.



Index Guidelines

Index Guidelines

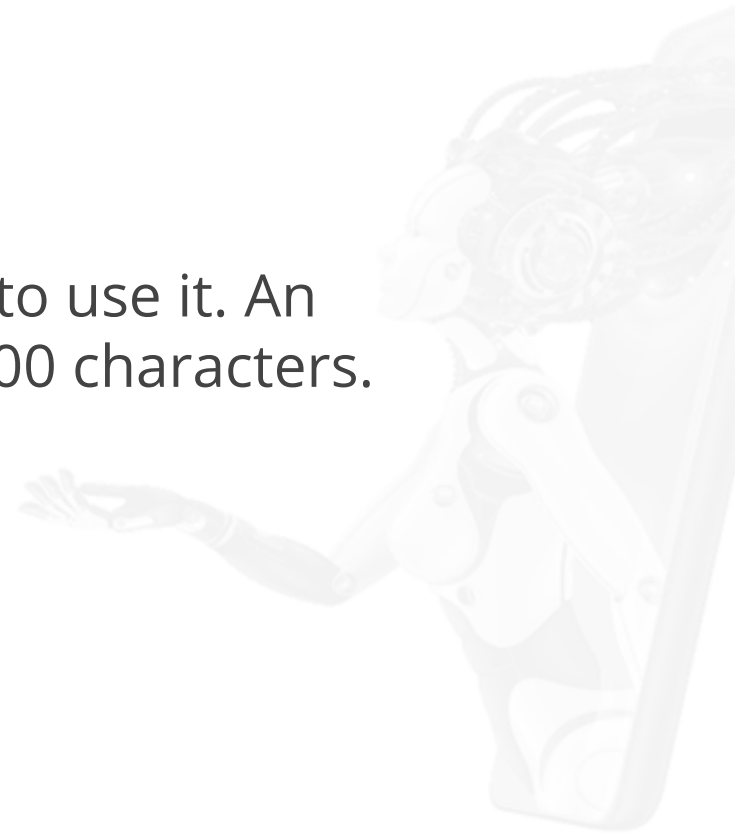
Choosing the right columns and types for an index is a crucial part of building a useful index.



Index Guidelines

Keep index keys as short as possible

The larger an index key gets, the harder it is for a database to use it. An integer key is smaller than a character field that can carry 100 characters. Keep clustered indexes as short as possible.



Index Guidelines


Keep separate index
keys

Indexes with a limited percentage of duplicated values are the most effective. With a decent index, the database will be able to ignore as many records as possible.

Index Guidelines

Keep selective indexes

A selective index has a lot of unique values. A unique index is the most selective of all the indexes, because there are no duplicate values.



Index Guidelines

Keeping indexes up-to-date

You will need to see existing indexes as well as delete or rename them, in addition to generating new ones. As the schema or even naming standards change, this is part of the database's continual maintenance cycle.

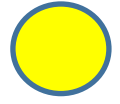
Creating a Clustered Index in SQL

Clustered Index in SQL

- A clustered index is an index that reorders the actual storage of entries in a table.
- Each table can only have one clustered index.

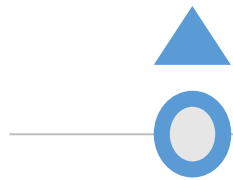


Clustered Index in SQL



Following are the essential characteristics of a clustered index:

It allows us to store both data and indexes at the same time.



It's an excellent choice for range or group queries that return min, max, or count values.



It only has one manner of storing data, which is dependent on the key values.



It always takes one or more columns to create an index.




Covering Queries With Indexes in SQL




A covered query is a query where all the columns in the query's result set are pulled from non-clustered indexes.



The careful placement of indexes transforms a query into a covered query.



Clustered Index in SQL



Problem Scenario:

A data analyst wants to sort a table in order, with the help of a clustered index. Create a primary key which acts as a clustered index in that table.

Objective:

Implement the clustered index to obtain the result.

Instructions:

Refer the emp_data table created before and perform the objectives.

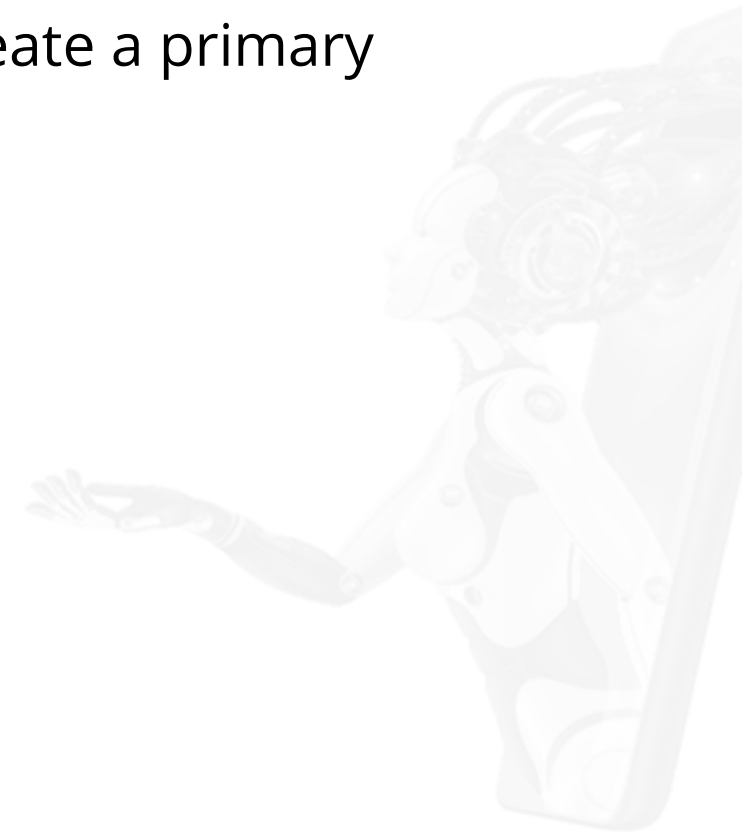


Table Description

Field Name	Description
EMP_ID	Employee ID
FIRST_NAME	First name of the employee
LAST_NAME	Last name of the employee
GENDER	Gender of the employee (M/F)
ROLE	Designation of the employee (Junior, Senior, Lead, and Associate Data Scientist)
DEPT	Name of the department (Retail, Finance, Automotive, and Healthcare)

Table Description

Field Name	Description
EXP	Experience of the employee
COUNTRY	Country where the employee lives
CONTINENT	Continent based on the country
SALARY	Salary of the employee per month
EMP_RATING	Rating for the employee (1: Not Achieving Any Goals, 2: Below Expectations, 3: Meeting Expectations, 4: Excellent Performance, 5: Overachiever)
MANAGER_ID	It is the employee ID for the manager

Clustered Index in SQL





In the query below, the PRIMARY KEY is a clustered index.

QUERY

```
CREATE TABLE  
emp_data( `EMP_ID` Varchar NOT NULL, `FIRST_NAME` varchar(45) DEFAULT NULL, `LAST_NAME` var  
char(3) DEFAULT NULL, `GENDER` varchar(20) DEFAULT NULL, `ROLE` varchar(25) DEFAULT NULL, `D  
EPT` varchar(25) DEFAULT NULL, `EXP` varchar(25) DEFAULT NULL, `COUNTRY` varchar(25) DEFAULT N  
ULL, `CONTINENT` varchar(25) DEFAULT NULL, `SALARY` INT DEFAULT NULL, `EMP_RATING` INT  
DEFAULT NULL, `MANAGER_ID` varchar(25) DEFAULT NULL, PRIMARY KEY (`EMP_ID`))  
clustered index UNIQUE KEY `SALARY` (`SALARY`);
```

Clustered Index in SQL

Output

Result Grid   Filter Rows: <input type="text"/> Export:  Wrap Cell Content: 											
	EMP_ID	FIRST_NAME	LAST_NAME	GENDER	ROLE	DEPT	EXP	COUNTRY	CONTINENT	SALARY	EMP_RATIN
▶	E260	Roy	Collins	M	SENIOR DATA SCIENTIST	RETAIL	7	INDIA	ASIA	7000	3
	E245	Nian	Zhen	M	SENIOR DATA SCIENTIST	RETAIL	6	CHINA	ASIA	6500	2
	E620	Katrina	Allen	F	JUNIOR DATA SCIENTIST	RETAIL	2	INDIA	ASIA	3000	1
	E640	Jenifer	Jhones	F	JUNIOR DATA SCIENTIST	RETAIL	1	COLOMBIA	SOUTH AMERICA	2800	4
	E403	Steve	Hoffman	M	ASSOCIATE DATA SCIENTIST	FINANCE	4	USA	NORTH AMERICA	5000	3
	E204	Karene	Nowak	F	SENIOR DATA SCIENTIST	AUTOMOTIVE	8	GERMANY	EUROPE	7500	5
	E204	Karene	Nowak	F	SENIOR DATA SCIENTIST	AUTOMOTIVE	8	GERMANY	EUROPE	7500	5
	E010	William	Butler	M	LEAD DATA SCIENTIST	AUTOMOTIVE	12	FRANCE	EUROPE	9000	2
	E478	David	Smith	M	ASSOCIATE DATA SCIENTIST	RETAIL	3	COLOMBIA	SOUTH AMERICA	4000	4
	E005	Eric	Hoffman	M	LEAD DATA SCIENTIST	FINANCE	11	USA	NORTH AMERICA	8500	3
	E532	Claire	Brennan	F	ASSOCIATE DATA SCIENTIST	AUTOMOTIVE	3	GERMANY	EUROPE	4300	1
	E583	Janet	Hale	F	MANAGER	RETAIL	14	COLOMBIA	SOUTH AMERICA	10000	2
	E103	Emily	Grove	F	MANAGER	FINANCE	14	CANADA	NORTH AMERICA	10500	4
	E612	Tracy	Norris	F	MANAGER	RETAIL	13	INDIA	ASIA	8500	4
	E428	Pete	Allen	M	MANAGER	AUTOMOTIVE	14	GERMANY	EUROPE	11000	4
	E002	Cynthia	Brooks	F	PRESIDENT	ALL	17	CANADA	NORTH AMERICA	14500	5
	E002	Cynthia	Brooks	F	PRESIDENT	ALL	17	CANADA	NORTH AMERICA	14500	5

Assisted Practice: Covering Query



Duration: 15 min

Problem Statement: You are required to create the temperature as unique, add new index columns, and view the index of the weather table for data integrity.

ASSISTED PRACTICE

Assisted Practice: Covering Query



Steps to be performed:

Step 1 : Creating the weather table and inserting values in it:

CREATE

```
CREATE TABLE weather (    temp INT NOT NULL,    windspeed  
varchar(45) NOT NULL,    vapour varchar(45) NOT NULL,  
climate varchar(45) NOT NULL );
```

INSERT

```
INSERT INTO weather(temp,windspeed,vapour,climate) VALUES  
( '35', '120km/hr', '21', 'summer' );
```


Assisted Practice: Covering Query



Step 2 : Querying to create unique index:

QUERY

```
CREATE UNIQUE INDEX temp_index ON weather(temp);
```

Step 3 : Querying to add new index:

QUERY

```
ALTER TABLE weather ADD INDEX  
index_col1_col2(windspeed,vapour,climate);
```

ASSISTED PRACTICE

Assisted Practice: Covering Query



Step 3 : Querying to show index:

QUERY

```
SHOW INDEX FROM weather;
```

Output :

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
weather	0	temp_index	1	temp	A	0	NULL	NULL		BTREE			YES	NULL
weather	1	index_co1_col2	1	windspeed	A	6	NULL	NULL		BTREE			YES	NULL
weather	1	index_co1_col2	2	vapour	A	6	NULL	NULL		BTREE			YES	NULL
weather	1	index_co1_col2	3	climate	A	6	NULL	NULL		BTREE			YES	NULL

ASSISTED PRACTICE


Common Table Expression

Common Table Expression

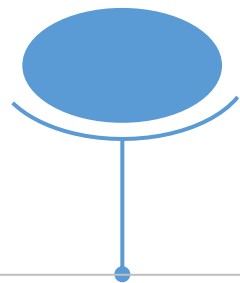


Each statement or query in MySQL generates a temporary result or relation. CREATE, INSERT, SELECT, UPDATE, DELETE, and other statements employ a common table expression (CTE) to name the temporary results sets that exist within the execution scope of that statement.

Common Table Expression





Following are some of the most important aspects of CTE:




WITH clause is used to define CTE.

WITH clause allows you to define several CTEs in a single query.



CTE makes code maintainability easier.

The execution scope of CTE exists within the statement in which it is used.



Common Table Expression

The basic syntax of CTE in MySQL is as follows:

SYNTAX

```
WITH cte_name (column_names) AS (query)  
SELECT * FROM cte_name;
```



Common Table Expression

Problem Scenario:

A data analyst wants to apply a CTE to the table with first name, last name, and the country of the person whose salary is greater than or equal to 4000.

Objective:

Implement the CTE to obtain the required result.

Instructions:

Refer the emp_data table which was created and shown before. Perform the objectives mentioned above.



Common Table Expression

Following **emp_data** table is used to show the CTE:

Result Grid

Filter Rows:

Export:

Wrap Cell Content:

	EMP_ID	FIRST_NAME	LAST_NAME	GENDER	ROLE	DEPT	EXP	COUNTRY	CONTINENT	SALARY	EMP_RATING	MANAGER_ID
	E204	Karene	Nowak	F	SENIOR DATA SCIENTIST	AUTOMOTIVE	8	GERMANY	EUROPE	7500	5	E428
	E428	Pete	Allen	M	MANAGER	AUTOMOTIVE	14	GERMANY	GERMANY	11000	4	E002
▶	E478	David	Smith	M	ASSOCIATE DATA SCIENTIST	RETAIL	3	COLOMBIA	SOUTH AMERICA	4000	4	E583
	E532	Claire	Brennan	F	ASSOCIATE DATA SCIENTIST	AUTOMOTIVE	3	GERMANY	EUROPE	4300	1	E428
	E583	Janet	Hale	F	MANAGER	RETAIL	14	COLOMBIA	SOUTH AMERICA	10000	2	E002
	E612	Tracy	Norris	F	MANAGER	RETAIL	13	INDIA	ASIA	8500	4	E002
▼	E204	Karene	Nowak	F	SENIOR DATA SCIENTIST	AUTOMOTIVE	8	GERMANY	EUROPE	7500	5	E428
< >	E245	Nian	Zhen	M	SENIOR DATA SCIENTIST	RETAIL	6	CHINA	ASIA	6500	2	E583
↶ ↷	E260	Roy	Collins	M	SENIOR DATA SCIENTIST	RETAIL	7	INDIA	ASIA	7000	3	E583
Informa	E403	Steve	Hoffman	M	ASSOCIATE DATA SCIENTIST	FINANCE	4	USA	NORTH AMERICA	5000	3	E103
	E620	Katrina	Allen	F	JUNIOR DATA SCIENTIST	RETAIL	2	INDIA	ASIA	3000	1	E612
▲	E640	Jenifer	Jhones	F	JUNIOR DATA SCIENTIST	RETAIL	1	COLOMBIA	SOUTH AMERICA	2800	4	E612

emp_data 3 x

Common Table Expression

CTE QUERY

```
WITH emp_in_Germany AS  
(  
  SELECT * FROM sys . emp_data WHERE COUNTRY = 'GERMANY'  
)  
  
SELECT FIRST_NAME , SALARY , COUNTRY from emp_in_Germany  WHERE SALARY >= '4000' order by  FIRST_NAME;
```

Common Table Expression

Output:

<

Result Grid

Filter Rows:

Export:

Wrap Cell Content:

	FIRST_NAME	SALARY	COUNTRY
▶	Karene	7500	GERMANY
	Karene	7500	GERMANY
	Claire	4300	GERMANY

SQL Best Practices

SQL Best Practices

Always check for NULLS in your data

Don't use a query again if it does not serve its entire function

Always keep an eye on the execution plan and track of the time costs

Avoid sub queries, and do joins or write functions if required

SQL Best Practices

SQL Best Practices

Use the right indexes for faster search results

Improve readability and maintainability, and ensure that the correct columns are retrieved

When your SQL statement has several sources, always utilize table aliases

In your INSERT statements, always use a column list



SQL Best Practices

Key Takeaways

- Clustered index always uses one or more columns for creating an index.
- Choosing the right columns and types for an index is a crucial part of building a useful index.
- A covered query is a query where all the columns in the query's result set are pulled from non-clustered indexes.
- XML indexes can be constructed on XML data type columns. They boost query performance by indexing all tags, values, and routes across XML instances in the column.





Knowledge Check

**Knowledge
Check**

1

Which of the following is one of the best practices in SQL?

- A. Use the right indexes for slower search results
- B. Make subqueries, and do not make joins or write functions
- C. Always check for NULLS in your data
- D. Use a query again if it does not serve its entire function



Knowledge Check

1

Which of the following is one of the best practices in SQL?

- A. Use the right indexes for slower search results
- B. Make subqueries, and do not make joins or write functions
- C. Always check for NULLS in your data
- D. Use a query again if it does not serve its entire function



The correct answer is: **C**

Always check for NULLS in your data. This is one of the main SQL practices.

**Knowledge
Check**

2

Which of the following stores Unicode characters?

- A. CHAR
- B. VARCHAR
- C. NVARCHAR
- D. None of the above



Knowledge
Check

2

Which of the following stores Unicode characters?

- A. CHAR
- B. VARCHAR
- C. NVARCHAR
- D. None of the above



The correct answer is: **C**

Unicode characters are stored in NVARCHAR.

**Knowledge
Check**

3

Which of the following are index guidelines?

- A. Keep index keys as short as possible
- B. Keep separate index keys
- C. Not keeping indexes up-to-date
- D. Keeping a single key for indexes



Knowledge
Check

3

Which of the following are index guidelines?

- A. Keep index keys as short as possible
- B. Keep separate index keys
- C. Not keeping indexes up-to-date
- D. Keeping a single key for indexes



The correct answers are: **A,B**

Index guidelines are keeping separate index keys and making them as short as possible.